

UNIVERSITY OF KENT

A systematic analysis of attacks on a honeypot

CO880 - Project and Dissertation

Yann Ferrere
ID 15906875
yf57@kent.ac.uk

ABSTRACT

ACKNOWLEDGEMENT

Table of contents

1	Introduction	5
1.1	Context	5
1.2	Goals of the project	5
1.3	Contribution	5
1.4	Structure of the disseration	6
2	Literature review and related work	6
3	Methodology	6
4	Implementation and configuration	6
4.1	Honeypots	6
4.1.1	Honeypots overview	6
4.1.2	Cowrie and its configuration	7
4.1.3	Kippograph and data vizualisation	8
4.2	Malware string analyzer	8
4.2.1	Strings categories	8
4.2.2	Strings extraction	9
4.2.3	Virus total analysis	9
5	Analysis and results	9
5.1	Attacks overview	9
5.2	Authentication attempts	9
5.3	Malicious downloads	9
5.4	Malware binaries analysis	9
6	Future research	10
7	Concluding discussion	10
8	Bibliography	10
9	Appendices	10

LIST OF FIGURES

1 Introduction

1.1 Context

With the development of the Internet, the number of computer servers has significantly increased. Indeed, Internet is based on interconnection between all devices such as computers and smartphone but also servers. The aim of those servers is to store data such as video, image, music, databases, but it also provides interaction between software executed on them and remote users. These interactions are called client/server communication. Servers can, for instance, execute a software such as a web server, which consists of making accessible web pages stored on it to remote users. Evidently, plenty of other servers usages exists but all of them widely implies to store data, run software and communicate information on the network.

Meanwhile, the usage of devices connected to internet has become extremely widespread. Some basic tasks of daily living such as consulting its bank accounts, communicate with each other or read the news, can be performed anywhere by using an internet connected device such as a smartphone. All of these services can be provided by a private company, but also by a government or an individual. In order to function efficiently, these services are developed by using a computer language and are hosted on a server to store users data and to be accessible remotely.

However, all of these communication and stored data can be more or fewer sensitives. Indeed, communication between two people, bank account information, password, and other credentials data have a value that is interesting to steal or alter by an ill-intentioned person.

1.2 Goals of the project

The risk of a server attack is now becoming ubiquitous. Based on this known risk I decided to analyze what are the steps of an attack on a server. Moreover, servers work as common computers and need to have an operating system to operate properly. This dissertation will focus on Linux servers distributions. Indeed, the majority of servers are Linux based and can run, for instance, known OS such as Ubuntu or Debian.

Consequently, this project consists of putting in place a Linux server accessible remotely by its unique public IP address and then to install a Honeypot on it. This type of tool is a software that emulates a fake Linux environment in order to analyze attackers behaviour. Moreover, a honeypot makes a server accessible remotely through a fake SSH connection with really common identifiers (username/password) in order to be bruteforced easily. Hence, this tool provides a way to allow malicious connection on its server while restricting the access to the machine.

During 6 months, our honeypot server received around 72500 SSH login attempts with 62 percent of successful connection. Those attacks were from 2550 distinct IP addresses and from a large number of different countries all around the world. Following these attacks multiple data such as command lines executed by attackers, scripts and malware downloaded have allowed me to perform a more detailed analysis of attackers behaviours and their motivations.

1.3 Contribution

The honeypot technic to collect attacks on a Linux server is not new. Several papers have already put in place this configuration in order to analyze attackers behaviours. However, there are plenty of ways to implement a honeypot. Some papers have created their own network architectures with multiple machines, while others have used existing solutions such as Kippo. In this dissertation, the Honeypot is cowrie, a fork of the kippo project which is not longer maintained. Consequently, this paper shows in a real case scenario that this honeypot can be an appropriate choice in a project that needs to collect data from Linux server attacks.

It is also important to note that attackers technics to compromise Linux server are constantly evolving. Indeed, security researchers and developers are continuing to look at what are the newest vulnerabilities and try to patch them (CVE). Then, attackers find new attacks scenario in order to bypass new security detection and protection tools and so on. However, security researchers have to put sensors in order to collect data and analyze security issues. Consequently, all along this project research, a databases containing all attacks information has been created in order to update and improve knowledge of current Linux server attacks.

Furthermore, because of a large number of ELF malicious collected binaries I performed a static analysis of them by analyzing strings. In fact, when malware are not packed or stripped multiple hard coded strings can be extracted from them. During the analysis of these strings, 10 different types were observed such as IP addresses, URLs or readable english messages. In order to automate their extraction, I developed a python script that automatically extract strings from a given malware, extract, sort and store them into a database. The aim of this could be to generate a large dataset from strings contained in malware to be applied through a data mining algorithm in order to perform a pre-analysis of a malware.

1.4 Structure of the disertation

In order to understand the functioning of Honeypots and their objectives, a short review of the existing literature about similar project already performed is exposed first. Then, the methodology of my research will be explained.

The technical implementation and configuration of the server, the honeypot and tools used to store and read collected data will be examined in details. In the same way, the development and usage of the malware strings analyzer will be shown.

A concluding discussion will summarize this analysis and results will be described in the same section. Finally, in a future research part will provide a way to a possible future development and reflexion.

2 Literature review and related work

3 Methodology

4 Implementation and configuration

4.1 Honeypots

As we have seen previously, the analysis of Linux servers attacks requires collecting the maximum of information about each of them. In order to perform this task, I choose to put in place a honeypot software called Cowrie. In this section will be explained how works a honeypot, how cowrie has to be configured and finally how it is possible to extract useful attacks information from all collected data.

4.1.1 Honeypots overview

As seen during the literature and related works review, Honeypots are not a recent technics and are still commonly used to trap attackers. Moreover, they can be used in multiple manners to analyse different types of attacks. However, all of them are aimed at being undetectable by attackers. Indeed, honeypots software are built to receive attacks from real attackers who wants to compromise a vulnerable machine and confine the attack to a controlled environment. Hence, the configuration part of this type of tool and its security has to be constantly monitored. A honeypot containing an unknown vulnerability (also called 0day) or with incorrect settings, could be exploited by attackers and consequently lead the machine to be compromised.

Multiple systems can be used and some honeypots project, for instance, are able to be installed on smartphones by using a tool such as HoneyDroid. Moreover, the HoneyNet Project gather security specialist that use multiple types of honeypots in order to collect data and report their knowledge about detected attacks on all platforms. They also developed various and sundry tools such as HoneySnap, a script used to extract malicious events from large logs generated by their honeypots, and allowing to understand attacker's methodology on the inside of the compromised system.

Honeypots can also be associated together and create HoneyNets. Unlike a honeypot that can be installed on a system in production inside a complete network as a sensor or alone on a separate machine, a honeyNet is composed

of multiple honeypots. Each of them can allow to detect and analyse more complex attacks in a large network closer to reality for a company than a single machine.

In order to be attacked a honeypot has to be vulnerable and accessible remotely. There are a large number of possible vulnerabilities depending on the system targeted. Each of them can allow attackers to execute more or less critical tasks such as obtained a root access to a server. A known honeypots technics, to perform this tasks, is to emulate a fake Secure Shell (SSH) access. This service is used to obtain a remote secure shell from a server, in order to monitor it. Indeed, this remote shell can be accessed by knowing a username and its associated password, available on the server. Once the honeypot emulates a fake SSH access it will allow creating multiple valid identifiers allowing an attacker to access to a fake Linux environment. The objective is to set a very common username and password to facilitate the access to an attacker.

4.1.2 Cowrie and its configuration

Cowrie is an SSH honeypot allowing its user to log brute force attacks and the entire shell command lines performed by attackers once logged on the server. Developed in Python by Michel Oosterhof, this project is a fork of another well known SSH honeypot called Kippo and developed by desaster. Indeed, the developer of Cowrie decided to add many features such as the addition of the "ssh exec commands" support. Indeed, it is possible to execute command lines through SSH after being connected or execute a single command line remotely, without obtaining a full shell but only the output of its command. This is commonly used by attackers that developed scripts to automate their attacks and is consequently useful to be implemented on an SSH honeypot. Moreover, because the attackers who access to cowrie is in a fake Linux environment jailed from the rest of the real server operating systems, all command lines have to be re-implemented. Consequently, Kippo and then Cowrie honeypots have implemented commons command lines used by attackers and more generally all Linux users. However, one of them called "ping", used to test the communication between two devices on a network allowed in Kippo to ping IP addresses such as "555.555.555.555" which is an impossible IP address. Indeed, an IP address is composed of 4 bytes and bytes can only store a value between 0 and 255. Hence, Kippo project is no longer maintained and some automate scripts used to detect this honeypot have been developed.

Consequently, I decided to use this SSH honeypot in order to collect attacks data. Its installation is really basic. First of all, a server accessible by a public IP address is needed. For my research, I choose to subscribe to a Virtual Private Server (VPS) on OVH, a French servers provider. Unlike a dedicated server, a VPS is basically a virtual machine running on a physical machine with a chosen operating systems. Its advantage is its low cost, because of the multiple virtual machines that can be run at the same time on a single dedicated server, associated with a sufficient configuration for a honeypot (2gigabytes of RAM and 10 gigabytes of storage).

The second part consists of deploying the honeypot on this VPS. Cowrie is available on Github, a website allowing to store developers projects on a private or public repository with versions numbers. Hence, it just needed to install with the packet manager available on the chosen operating system the "git" command line in order to clone the cowrie repository. This previous step implies to have created a user "cowrie" on your system, only use to run the honeypot. Then, we have to install python and the associated libraries listed in the requirements.txt file. Once it has been done, the configuration file has to be edited to enable the MySQL database storage of all data collected by the honeypot and also set the port number of the honeypot to 22. Indeed, the SSH service is by default set in port 22, consequently, the server has to be configured to ensure that the SSH honeypot run on port 22 and the original SSH server on a chosen port. This step is essential to be sure to maintain an SSH access to monitor your server.

Before running this honeypot we can also define a list containing valid usernames and passwords to be connected to the server. Then, a shell script called "start.sh" has to be executed to run the honeypot. More details about the installation and configuration can be found easily on internet.

Finally, one important point is to ensure that the server use does not contain any sensitives information. Indeed, this is a software and despite the availability of the code on GitHub the authors cannot ensure that any vulnerabilities will be found.

4.1.3 Kippograph and data vizualisation

An SSH honeypot generate a large amount of data. These data may refer to all information related to SSH connection attempts such as identifiers used and IP address used, but also to all events when the attackers access to a shell on our server. In the case of cowrie and Kippo Honeypots, all data are stored on a MySQL database automatically. This database is structured in 8 tables (auth, clients, downloads, input, keyfingerprints, sensors, sessions and ttylog). However, due to the quantity of attacks that can be received on an SSH honeypots, it quickly became apparent that read information from a database were a difficult task. Consequently, I decided to put in place a graphical interface allowing to extract and format useful data contains on this database.

Moreover, its important to notice that Cowrie apply the same database structure as Kippo. Hence, due to the large community of Kippo, I have been able to use Kippograph (v1.5.1), a tool from their toolbox. This tool is a web interface using the Libchart PHP chart drawing library that gathers multiple libraries in order to build maps and graphs from given data. Consequently, this tool allowed me to perform a daily analysis of the actual state of attacks performed against my server. Indeed, a simple GET request to one of the 6 available pages (kippo-graph, kippo-input, kippo-playlog, kippo-ip, kippo-geo and graph gallery), return several graphs and information such as the actual number of login attempts and distinct source IP addresses, easy to understand and interpret.

The installation is also simple and fast to implement. First of all, Kippograph is a web interface consequently a webserver such as Apache or Nginx has to be installed on the server. Once, it has been installed all of the source code, located on GitHub has to be downloaded on the web server localization (by default /var/www/html). Then, the configuration file has to be edited to allow it to get data from the cowrie MySQL database. Finally, this website has to be protected with all basic security rules such as a strong password to avoid everyone to access to this information remotely. More details about the installation and configuration can be found easily on internet.

4.2 Malware string analyzer

By putting in place an SSH honeypot, I have been able to collect a large number of binary downloaded by attackers when there were connected to the server. Those executables and Linkable Format binaries, also called ELF binaries, contains hard coded strings on them. Based on these strings I decided to implement a python script, called Malware Strings Analyser (MSA), used to extract and sort strings into categories. This allowed me to automate this task to all gathered binaries and consequently to focus on the analysis of data generated by this script. The source code and documentation are available on GitHub.

4.2.1 Strings categories

The analysis of readable strings is useful in order to understand how works a malware without running it, also called a static analysis. However, this could be done only if binaries are not packed. Indeed, some malware developers can employ encryption algorithms in order to hide the active part of the malware to bypass possible detections of them. Moreover, if a binary is stripped then we could not be able to get readable strings such as the name of variables and functions symbols. Indeed, ELF binaries are commonly compiled with debug information. However, developers can use compilers options in order to remove all debug information in order to hide a maximum of information about their functionalities (included libraries, functions used,...).

During my research, I divided strings that I found in 10 categories:

1. IP addresses: it is possible to find hard coded IP addresses, with sometimes a port number (e.g. 93.174.89.143:23)
2. Identifiers: some binaries contains strings which look like username/passwords, probably to perform a password guessing brute force attack on the compromised server.
3. Command lines: there are some strings that contain readable shell command line that could be executed by the malware.
4. Url and files: it is possible to find strings that contain URL allowing to download a file, probably to update the malware remotely. Those URLs are often written into a command line string.
5. Path: some binaries contains strings with the absolute path to a Linux binary or a file (e.g. /bin/sh)

6. Symbols: when a binary is not stripped it is possible to get more specific information such as the name of used functions.
7. Format string: this type of string consists of building string dynamically (e.g. %d.%d.%d.%d for an IP address).
8. Display message: some readable message contains error, success or helps message. Those string can help to understand functionalities of the malware.
9. Sections: it is possible to get name of the program sections (e.g. .text, .data, .bss, ...).
10. Other: Some other strings contain readable information but I was not able to determine if they are useful and where they come from. This category also contains some extracted strings with random character without sense.

Consequently, Malware String Analyzer tool, extract all readable character contains in a binary file, sort them depending on their groups and store them in an SQLite database. Finally, it is important to notice that a malware's developer can also introduce incorrect readable strings in order to hide the correct functioning of its malware.

4.2.2 Strings extraction

4.2.3 Virus total analysis

5 Analysis and results

5.1 Attacks overview

5.2 Authentication attempts

5.3 Malicious downloads

5.4 Malware binaries analysis

- 6 Future research
- 7 Concluding discussion
- 8 Bibliography
- 9 Appendices